

# CS441\_SP24\_HW1\_Starter

February 5, 2024

## 0.1 CS441: Applied ML - HW 1

### 0.1.1 Parts 1-2: MNIST

Include all the code for generating MNIST results below

```
[2]: # initialization code
import numpy as np
from keras.datasets import mnist
%matplotlib inline
from matplotlib import pyplot as plt
from scipy import stats
import faiss
import time

def load_mnist():
    '''
    Loads, reshapes, and normalizes the data
    '''
    (x_train, y_train), (x_test, y_test) = mnist.load_data() # loads MNIST data
    x_train = np.reshape(x_train, (len(x_train), 28*28)) # reformat to 768-d
    ↪vectors
    x_test = np.reshape(x_test, (len(x_test), 28*28))
    maxval = x_train.max()
    x_train = x_train/maxval # normalize values to range from 0 to 1
    x_test = x_test/maxval
    return (x_train, y_train), (x_test, y_test)

def display_mnist(x, subplot_rows=1, subplot_cols=1):
    '''
    Displays one or more examples in a row or a grid
    '''
    if subplot_rows>1 or subplot_cols>1:
        fig, ax = plt.subplots(subplot_rows, subplot_cols, figsize=(15,15))
        for i in np.arange(len(x)):
            ax[i].imshow(np.reshape(x[i], (28,28)), cmap='gray')
            ax[i].axis('off')
    else:
```

```
plt.imshow(np.reshape(x, (28,28)), cmap='gray')
plt.axis('off')
plt.show()
```

```
2024-02-05 14:55:46.910061: E
external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:9261] Unable to register
cuDNN factory: Attempting to register factory for plugin cuDNN when one has
already been registered
2024-02-05 14:55:46.910141: E
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:607] Unable to register
cuFFT factory: Attempting to register factory for plugin cuFFT when one has
already been registered
2024-02-05 14:55:46.912560: E
external/local_xla/xla/stream_executor/cuda/cuda_blas.cc:1515] Unable to
register cuBLAS factory: Attempting to register factory for plugin cuBLAS when
one has already been registered
2024-02-05 14:55:46.922567: I tensorflow/core/platform/cpu_feature_guard.cc:182]
This TensorFlow binary is optimized to use available CPU instructions in
performance-critical operations.
To enable the following instructions: AVX2 FMA, in other operations, rebuild
TensorFlow with the appropriate compiler flags.
2024-02-05 14:55:47.603859: W
tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Could not
find TensorRT
```

```
[3]: # example of using MNIST load and display functions
(x_train, y_train), (x_test, y_test) = load_mnist()
display_mnist(x_train[:10], 1, 10)
print('Total size: train={}, test {}'.format(len(x_train), len(x_test)))
```



Total size: train=60000, test =10000

```
[3]: display_mnist(x_test[:10], 1, 10)
```



```
[4]: x_train[0]
```

[illegible]

0.99215686, 0.99215686, 0.99215686, 0.77647059, 0.71372549,  
 0.96862745, 0.94509804, 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0.31372549, 0.61176471, 0.41960784, 0.99215686, 0.99215686,  
 0.80392157, 0.04313725, 0. , 0.16862745, 0.60392157,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0.05490196,  
 0.00392157, 0.60392157, 0.99215686, 0.35294118, 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0.54509804,  
 0.99215686, 0.74509804, 0.00784314, 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0.04313725, 0.74509804, 0.99215686,  
 0.2745098 , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0.1372549 , 0.94509804, 0.88235294, 0.62745098,  
 0.42352941, 0.00392157, 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0.31764706, 0.94117647, 0.99215686, 0.99215686, 0.46666667,  
 0.09803922, 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0.17647059,  
 0.72941176, 0.99215686, 0.99215686, 0.58823529, 0.10588235,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0.0627451 , 0.36470588,  
 0.98823529, 0.99215686, 0.73333333, 0. , 0. ,

0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0.97647059, 0.99215686,  
 0.97647059, 0.25098039, 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0.18039216, 0.50980392,  
 0.71764706, 0.99215686, 0.99215686, 0.81176471, 0.00784314,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0.15294118,  
 0.58039216, 0.89803922, 0.99215686, 0.99215686, 0.99215686,  
 0.98039216, 0.71372549, 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0.09411765, 0.44705882, 0.86666667, 0.99215686, 0.99215686,  
 0.99215686, 0.99215686, 0.78823529, 0.30588235, 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0.09019608, 0.25882353, 0.83529412, 0.99215686,  
 0.99215686, 0.99215686, 0.99215686, 0.77647059, 0.31764706,  
 0.00784314, 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0.07058824, 0.67058824, 0.85882353,  
 0.99215686, 0.99215686, 0.99215686, 0.99215686, 0.76470588,  
 0.31372549, 0.03529412, 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0.21568627, 0.6745098 ,  
 0.88627451, 0.99215686, 0.99215686, 0.99215686, 0.99215686,  
 0.95686275, 0.52156863, 0.04313725, 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0.53333333, 0.99215686, 0.99215686, 0.99215686,  
 0.83137255, 0.52941176, 0.51764706, 0.0627451 , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,  
 0. , 0. , 0. , 0. , 0. ,

[illegible]

## 1. Retrieval, Clustering, and NN Classification

```
[4]: # Retrieval

def get_nearest(X_query, X):
    ''' Return the index of the sample in X that is closest to X_query according
        to L2 distance '''
    dist = np.square(X_query - X).sum(axis=-1)
    return np.argmin(dist)

j = get_nearest(x_test[0], x_train)
print(j)

j = get_nearest(x_test[1], x_train)
print(j)
```

53843

28882

```
[5]: # K-means

def kmeans(X, K, niter=10):
    '''
        Starting with the first K samples in X as cluster centers, iteratively assign
        ↪ each
        point to the nearest cluster and compute the mean of each cluster.
        Input: X[i] is the ith sample, K is the number of clusters, niter is the
        ↪ number of iterations
        Output: K cluster centers
    '''
```

```
'''
# TO DO -- add code to display cluster centers at each iteration also
centers = X[:K].copy()
for i in range(1,1+niter):
    print(f"Iteration {i}")
    clusters = [[] for _ in range(K)]
    for x in X:
        idx = get_nearest(x, centers)
        clusters[idx].append(x)
    for c in range(K):
        centers[c] = np.stack(clusters[c]).mean(axis=0)
    display_mnist(centers, 1, K)

K=30
centers = kmeans(x_train[:1000], K)
```

Iteration 1

5 0 4 1 9 2 1 3 1 4 3 5 3 6 1 7 0 8 6 9 4 0 4 1 2 2 9 3 2 1

Iteration 2

5 0 4 1 9 2 1 3 1 4 3 5 3 6 1 7 0 8 6 9 4 0 4 1 2 0 9 3 2 1

Iteration 3

5 0 4 1 9 2 1 3 1 4 3 5 3 6 1 7 0 8 6 9 4 0 4 1 2 0 9 3 2 1

Iteration 4

5 0 4 1 9 2 1 3 1 4 3 5 3 6 1 7 0 8 6 9 4 0 4 1 2 0 9 3 2 1

Iteration 5

5 0 4 1 9 2 1 3 1 4 3 5 3 6 1 7 0 8 6 9 4 0 4 1 2 0 9 3 2 1

Iteration 6

509192331435361708694041207529

Iteration 7

509192331435361708694041207529

Iteration 8

509192331435361708694041207529

Iteration 9

509192331435361708694041207529

Iteration 10

509192331435361708694041207529

```
[6]: # 1-NN
errors = []
for i,x in enumerate(x_test[:100]):
    idx = get_nearest(x, x_train[:10000])
    errors.append(y_train[idx] != y_test[i])
print(np.mean(errors))
```

0.08

2. Make it fast

```
[11]: # install libraries you need for part 2
!apt install libomp-dev
!pip install faiss-cpu
import faiss
import time
```



```

E: Could not open lock file /var/lib/dpkg/lock-frontent - open (13:
Permission denied)
E: Unable to acquire the dpkg frontend lock (/var/lib/dpkg/lock-
frontend), are you root?
Defaulting to user installation because normal site-packages is not writeable
Collecting faiss-cpu
  Downloading
faiss_cpu-1.7.4-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (17.6
MB)
                                17.6/17.6 MB
21.9 MB/s eta 0:00:0000:0100:01
Installing collected packages: faiss-cpu
Successfully installed faiss-cpu-1.7.4

```

```

[6]: # retrieval

# TO DO (check that you're using FAISS correctly)
index = faiss.IndexFlatL2(x_train.shape[1])
index.add(x_train) # add the data
dist, idx = index.search(x_test[:2],1)
print(idx)

```

```

[[53843]
 [28882]]

```

```

[7]: import faiss
import numpy as np
import matplotlib.pyplot as plt

def kmeans_fast(X, K, niter=10):
    X = X.astype(np.float32)
    d = X.shape[1]
    cluster_centers = X[:K].copy()
    index = faiss.IndexFlatL2(d)

    rmse_list = []

    for i in range(niter):
        index.reset()
        index.add(cluster_centers)
        D, I = index.search(X, 1) # D is the squared distances
        rmse = np.sqrt(np.mean(D)) # Compute the RMSE from the squared
        ↪ distances
        rmse_list.append(rmse)

    for k in range(K):
        points_in_cluster = X[I.squeeze() == k]

```

```

        if len(points_in_cluster) > 0:
            cluster_centers[k] = np.mean(points_in_cluster, axis=0)


    return cluster_centers, rmse_list

# Assuming x_train is defined and is the correct input data
# You can now run the kmeans_fast function and plot the results
K = 10
centers, rmse = kmeans_fast(x_train, K, niter=20)
plt.plot(np.arange(len(rmse)), rmse, label='K=10')

K=30
centers, rmse = kmeans_fast(x_train, K, niter=20)
plt.plot(np.arange(len(rmse)), rmse, label='K=30')

K=100
centers, rmse = kmeans_fast(x_train, K, niter=20)
plt.plot(np.arange(len(rmse)), rmse, label='K=100')
plt.legend(), plt.ylabel('RMSE'), plt.xlabel('# iterations')
plt.show()

```



```

[8]: # 1-NN

nsample = [100, 1000, 10000, 60000]

acc_exact = []
timing_exact = []
timing_lsh = []
acc_lsh = []

for s in nsample:
    start = time.time()
    index = faiss.IndexFlatL2(x_train.shape[1])
    index.add(x_train[:s])
    dist, idx = index.search(x_test, 1)
    pred = y_train[idx.squeeze()]
    acc = (pred == y_test).mean()
    acc_exact.append(acc)
    timing_exact.append(time.time() - start)


for s in nsample:
    start = time.time()
    dim = x_train.shape[1]
    index = faiss.IndexLSH(dim, dim)
    index.add(x_train[:s])
    dist, idx = index.search(x_test, 1)
    pred = y_train[idx.squeeze()]
    acc = (pred == y_test).mean()
    acc_lsh.append(acc)
    timing_lsh.append(time.time() - start)


acc_exact = np.array(acc_exact)
acc_lsh = np.array(acc_lsh)

plt.semilogx(nsample, 1-acc_exact, label='exact')
plt.semilogx(nsample, 1-acc_lsh, label='lsh')
plt.legend(), plt.ylabel('% error'), plt.xlabel('# training samples')
plt.show()

plt.semilogx(nsample, timing_exact, label='exact')
plt.semilogx(nsample, timing_lsh, label='lsh')
plt.legend(), plt.ylabel('time'), plt.xlabel('# training samples')
plt.show()

```





```
[9]: # Confusion matrix
import sklearn
from sklearn.metrics import confusion_matrix

index = faiss.IndexFlatL2(x_train.shape[1])
index.add(x_train[:s])
dist, idx = index.search(x_test,1)
pred = y_train[idx.squeeze()]

C = confusion_matrix(y_test,pred)
print(C)
print(C[:, 2])

# TO DO
```

```
[[ 973   1   1   0   0   1   3   1   0   0]
 [   0 1129   3   0   1   1   1   0   0   0]
 [   7   6  992   5   1   0   2  16   3   0]
```

```
[ 0  1  2 970  1 19  0  7  7  3]
[ 0  7  0  0 944  0  3  5  1 22]
[ 1  1  0 12  2 860  5  1  6  4]
[ 4  2  0  0  3  5 944  0  0  0]
[ 0 14  6  2  4  0  0 992  0 10]
[ 6  1  3 14  5 13  3  4 920  5]
[ 2  5  1  6 10  5  1 11  1 967]]
[ 1  3 992  2  0  0  0  6  3  1]
```

## 0.2 Part 3: Temperature Regression

Include all your code used for part 2 in this section.

```
[7]: import numpy as np
# from google.colab import drive
%matplotlib inline
from matplotlib import pyplot as plt
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso

# load data (modify to match your data directory or comment)
def load_temp_data():
    # drive.mount('/content/drive')
    # datadir = "/content/drive/My Drive/CS441/24SP/hw1/"
    datadir = "./"
    T = np.load(datadir + 'temperature_data.npz')
    x_train, y_train, x_val, y_val, x_test, y_test, dates_train, dates_val, \
    ↪ dates_test, feature_to_city, feature_to_day = \
    T['x_train'], T['y_train'], T['x_val'], T['y_val'], T['x_test'], T['y_test'], \
    ↪ T['dates_train'], T['dates_val'], T['dates_test'], T['feature_to_city'], \
    ↪ T['feature_to_day']
    return (x_train, y_train, x_val, y_val, x_test, y_test, dates_train, \
    ↪ dates_val, dates_test, feature_to_city, feature_to_day)

# plot one data point for listed cities and target date
def plot_temps(x, y, cities, feature_to_city, feature_to_day, target_date):
    nc = len(cities)
    ndays = 5
    xplot = np.array([-5,-4,-3,-2,-1])
    yplot = np.zeros((nc,ndays))
    for f in np.arange(len(x)):
        for c in np.arange(nc):
            if cities[c]==feature_to_city[f]:
                yplot[feature_to_day[f]+ndays,c] = x[f]
    plt.plot(xplot,yplot)
    plt.legend(cities)
    plt.plot(0, y, 'b*', markersize=10)
    plt.title('Predict Temp for Cleveland on ' + target_date)
```


```
plt.xlabel('Day')
plt.ylabel('Avg Temp (C)')
plt.show()
```

```
[8]: # load data
(x_train, y_train, x_val, y_val, x_test, y_test, dates_train, dates_val,
  ↪ dates_test, feature_to_city, feature_to_day) = load_temp_data()
''' Data format:
    x_train, y_train: features and target value for each training sample
    ↪ (used to fit model)
    x_val, y_val: features and target value for each validation sample (used
    ↪ to select hyperparameters, such as regularization and K)
    x_test, y_test: features and target value for each test sample (used to
    ↪ evaluate final performance)
    dates_xxx: date of the target value for the corresponding sample
    feature_to_city: maps from a feature number to the city
    feature_to_day: maps from a feature number to a day relative to the
    ↪ target value, e.g. -2 means two days before
    Note: 361 is the temperature of Cleveland on the previous day
'''
f = 361
print('Feature {}: city = {}, day= {}'.format(f, feature_to_city[f],
  ↪ feature_to_day[f]))
baseline_rmse = np.sqrt(np.mean((y_val[1:]-y_val[:-1])**2)) # root mean squared
  ↪ error example
print('Baseline - prediction using previous day: RMSE={}'.format(baseline_rmse))


# plot first two x/y for val
plot_temps(x_val[0], y_val[0], ['Cleveland', 'New York', 'Chicago', 'Denver',
  ↪ 'St. Louis'], feature_to_city, feature_to_day, dates_val[0])
plot_temps(x_val[1], y_val[1], ['Cleveland', 'New York', 'Chicago', 'Denver',
  ↪ 'St. Louis'], feature_to_city, feature_to_day, dates_val[1])
```

Feature 361: city = Cleveland, day= -1

Baseline - prediction using previous day: RMSE=3.460601246750482







[9]: # K-NN Regression

```
def regress_KNN(X_trn, y_trn, X_tst, K=1):
    '''
    Predict the target value for each data point in X_tst using a
    K-nearest neighbor regressor based on (X_trn, y_trn), with L2 distance.
    Input: X_trn[i] is the ith training data. y_trn[i] is the ith training label.
    ↪ K is the number of closest neighbors to use.
    Output: return y_pred, where y_pred[i] is the predicted ith test value
    '''
    index = faiss.IndexFlatL2(X_trn.shape[1])
    index.add(X_trn)
    dist, idx = index.search(X_tst, K)
    pred = y_trn[idx]
    pred = pred.mean(axis=1)
    return pred

def normalize_features(x, y, fnum):
    ''' Normalize the features in x and y.
    For each data sample i:
    '''
```

```

        x2[i] = x[i]-x[i,fnum]
        y2[i] = y[i]-x[i,fnum]
    ...
    xnorm = x - np.expand_dims(x[:,fnum],1)
    ynorm = y - x[:,fnum]
    return xnorm,ynorm
# TO DO

k=5

# KNN with original features
y_pred = regress_KNN(x_train,y_train,x_test, K=k)
print(np.square(y_pred - y_test).mean()*0.5)
# TO DO

# KNN with normalized features
fnum = 361 # previous day temp in Cleveland

# TO DO
xnorm_train,ynorm_train = normalize_features(x_train,y_train,fnum)
xnorm_test, ynorm_test = normalize_features(x_test,y_test,fnum)

y_pred = regress_KNN(xnorm_train,ynorm_train,xnorm_test, K=k)

print(np.square(y_pred - ynorm_test).mean()*0.5)

```

3.249556245363484  
2.9324389176041588

### 0.3 Part 5: Stretch Goals

Include all your code used for part 5 in this section. You can copy-paste code from parts 1-3 if it is re-usable.

```

[3]: from scipy.stats import mode
# Stretch: KNN classification (Select K)
(x_train, y_train), (x_test, y_test) = load_mnist()

for K in [1, 3, 5, 11, 25]:
    index = faiss.IndexFlatL2(x_train.shape[1])
    index.add(x_train[:50000])
    dist, idx = index.search(x_train[50000:], K)
    pred = y_train[:50000][idx]
    mode_result, mode_count = mode(pred, axis=1)
    error = (mode_result != y_train[50000:]).mean()
    print(f"{K}: {error}")

```

```
1: 0.0288
3: 0.028
5: 0.0282
11: 0.0308
25: 0.0382
```

```
[5]: # test
K=3
index = faiss.IndexFlatL2(x_train.shape[1])
index.add(x_train)
dist, idx = index.search(x_test, K)
pred = y_train[idx]
mode_result, mode_count = mode(pred, axis=1)
error = (mode_result != y_test).mean()
print(f"{K}: {error}")
```

```
3: 0.0295
```

```
[14]: (x_train, y_train, x_val, y_val, x_test, y_test, dates_train, dates_val,
      ↪ dates_test, feature_to_city, feature_to_day) = load_temp_data()

for k in [1, 3, 5, 11, 25]:
    print(f"k: {k}")
    # KNN with original features
    y_pred = regress_KNN(x_train,y_train,x_val, K=k)
    print(np.square(y_pred - y_val).mean()*0.5)

    # KNN with normalized features
    fnum = 361 # previous day temp in Cleveland

    xnorm_train,ynorm_train = normalize_features(x_train,y_train,fnum)
    xnorm_val, ynorm_val = normalize_features(x_val,y_val,fnum)

    y_pred = regress_KNN(xnorm_train,ynorm_train,xnorm_val, K=k)

    print(np.square(y_pred - ynorm_val).mean()*0.5)

print("Optimal:")
xnorm_train,ynorm_train = normalize_features(x_train,y_train,fnum)
xnorm_test, ynorm_test = normalize_features(x_test,y_test,fnum)
y_pred = regress_KNN(xnorm_train,ynorm_train,xnorm_test, K=11)
print(np.square(y_pred - ynorm_test).mean()*0.5)
```

```
k: 1
4.330006390580244
3.8669564684344615
```

```

k: 3
3.2266843182278797
3.174005914679923
k: 5
3.095887258575192
3.032493865875963
k: 11
3.055782885507432
2.8908415810941466
k: 25
3.060219056497986
2.9103974575813334
Optimal:
2.7671311757775685

```

```

[28]: # Stretch: K-means (more iters vs redos)
(x_train, y_train), (x_test, y_test) = load_mnist()

ps = [(20,1),(4,5),(50,1),(10,5)]
for (ni,nr) in ps:
    print(ni,nr)
    res = []
    for i in range(5):
        kmeans = faiss.Kmeans(x_train.shape[1], 30, niter=ni, nredo=nr, seed=int(i))
        kmeans.train(x_train)
        dist, idx = kmeans.index.search(x_train, 1)
        rmse = np.sqrt(np.sum(dist) / x_train.shape[0])
        res.append(rmse)
    print(np.mean(res), np.std(res))
    print()

```

```

20 1
5.786271422463997 0.007644752654037744

4 5
5.822844075465655 0.012136287784838528

50 1
5.777090880990907 0.005488117925290951

10 5
5.787613203649412 0.0037177425483455743

```

```

[ ]: # from https://gist.github.com/jonathanagustin/b67b97ef12c53a8dec27b343dca4abba
# For use in Colab. For local, just use jupyter nbconvert directly

```

```

import os
# @title Convert Notebook to PDF. Save Notebook to given directory
NOTEBOOKS_DIR = "/content/drive/My Drive/CS441/24SP/hw1" # @param {type:
↳ "string"}
NOTEBOOK_NAME = "CS441_SP24_HW1_Solution.ipynb" # @param {type:"string"}
#-----#
from google.colab import drive
drive.mount("/content/drive/", force_remount=True)
NOTEBOOK_PATH = f"{NOTEBOOKS_DIR}/{NOTEBOOK_NAME}"
assert os.path.exists(NOTEBOOK_PATH), f"NOTEBOOK NOT FOUND: {NOTEBOOK_PATH}"
!apt install -y texlive-xetex texlive-fonts-recommended texlive-generic >
↳ /dev/null 2>&1
!jupyter nbconvert "$NOTEBOOK_PATH" --to pdf > /dev/null 2>&1
NOTEBOOK_PDF = NOTEBOOK_PATH.rsplit('.', 1)[0] + '.pdf'
assert os.path.exists(NOTEBOOK_PDF), f"ERROR MAKING PDF: {NOTEBOOK_PDF}"
print(f"PDF CREATED: {NOTEBOOK_PDF}")

```