

Name:

Matthew Tang

Netid:

mhtang2

## CS 441 - HW2: PCA and Linear Models

Complete the sections below. You do not need to fill out the checklist.

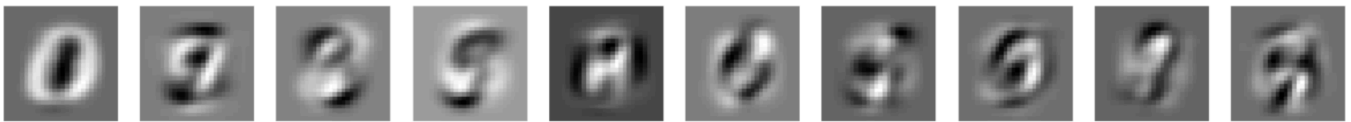
Total Points Available

[ ] / 160

1. PCA on MNIST
  - a. Display 10 principal component vectors [ ] / 5
  - b. Display scatterplot [ ] / 5
  - c. Plot cumulative explained variance [ ] / 5
  - d. Compression and 1-NN experiment [ ] / 15
2. MNIST Classification with Linear Models
  - a. LLR / SVM error vs training size [ ] / 20
  - b. Error visualization [ ] / 10
  - c. Parameter selection experiments [ ] / 15
3. Temperature Regression
  - a. Linear regression test [ ] / 10
  - b. Feature selection results [ ] / 15
4. Stretch Goals
  - a. PR and ROC curves [ ] / 10
  - b. Visualize weights [ ] / 10
  - c. Other embeddings [ ] / 15
  - d. One city is all you need [ ] / 15
  - e. SVM with RBF kernel [ ] / 10

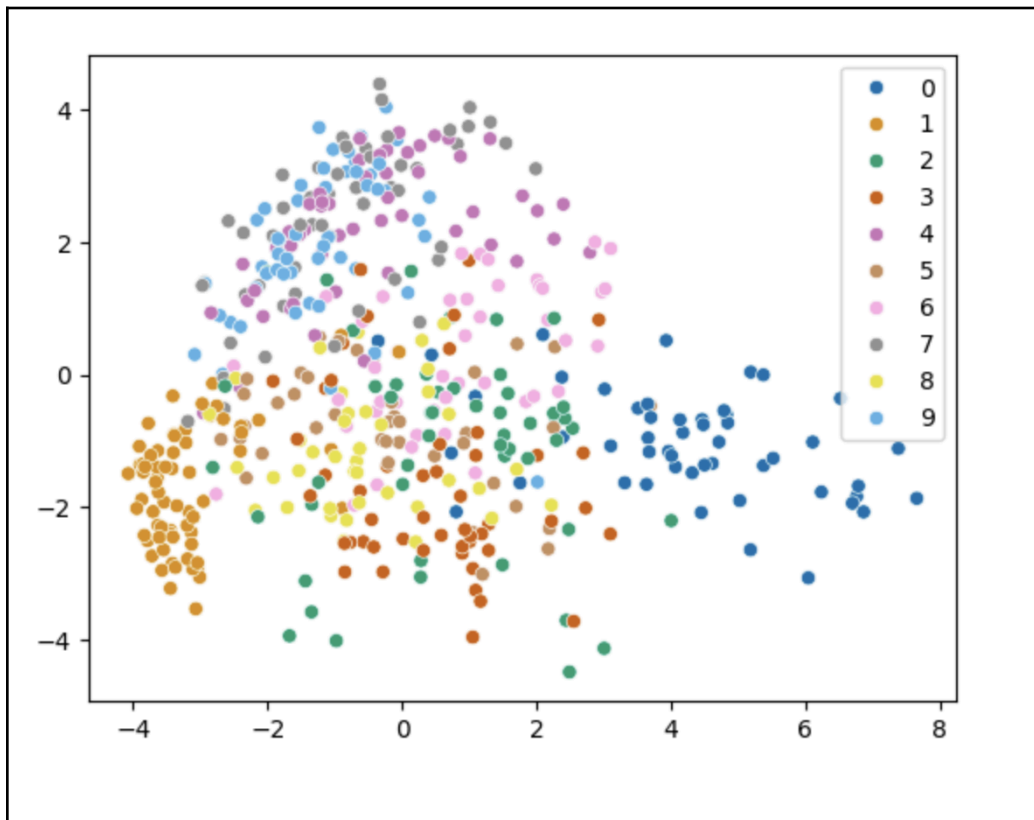
### 1. PCA on MNIST

#### a. Display 10 principal component vectors

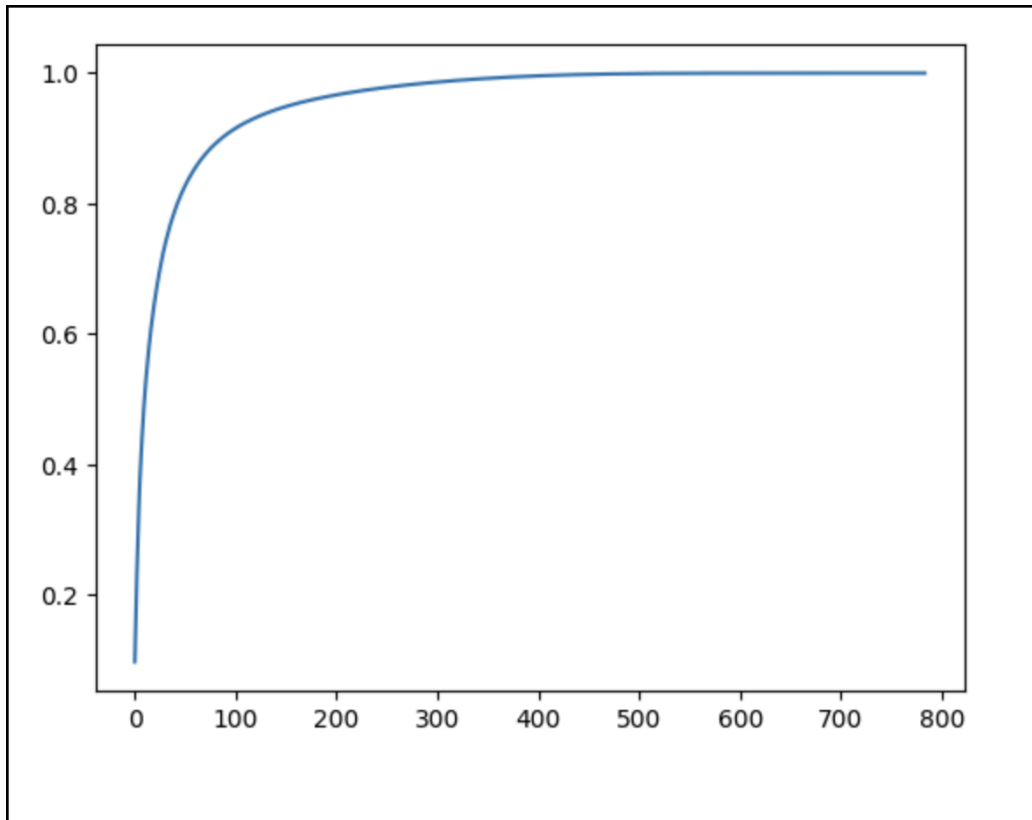


**b. Display scatterplot**

Scatterplot `x_train[:500]` for the first two PCA dimensions. Show a different color for each label.



**c. Plot cumulative explained variance**



#### d. Compression and 1-NN experiment

Number of components selected

	Total Time (s)	Test Error (%)	Dimensions
Brute Force (PCA)	4.545	2.73%	87
Brute Force	36.031	3.09%	784

## 2. MNIST Classification with Linear Models

#### a. LLR / SVM error vs training size

Test error (%)

# training samples	LLR	SVM
100	32.5%	32.4%
1,000	13.64 %	16.11 %

10,000	9.5 %	11.12 %
60,000	7.38 %	8.17 %

**b. Error visualization**

LLR



LLR Most confident correct:



LLR Most confident incorrect:



SVM

Most confident correct:



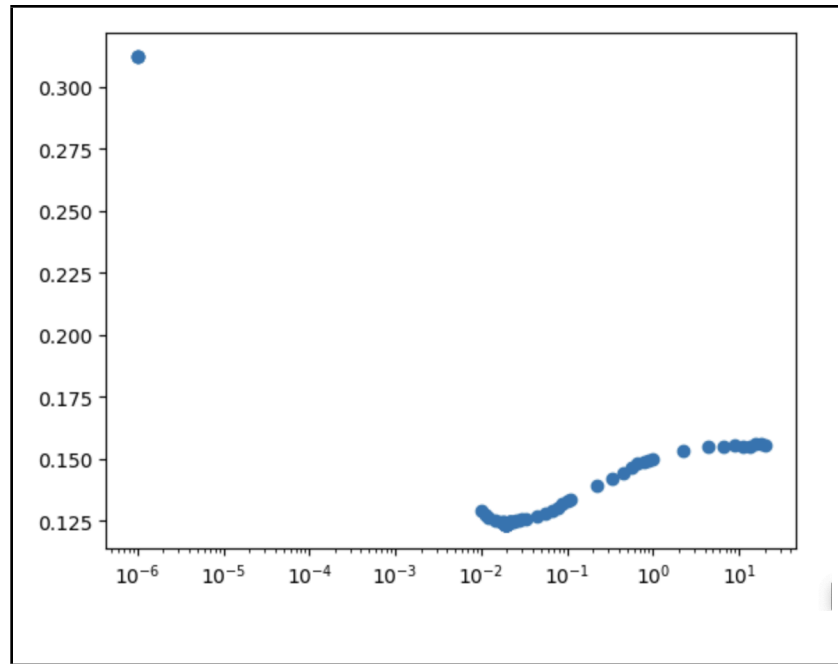
Most confident incorrect:



**c. Parameter selection experiments**

	SVM
Best C value	0.0196
Validation error (%)	12.35 %
Test error (%)	13.58 %

Plot C value vs validation error for values tested



### 3. Temperature Regression

#### a. Linear regression test

Test RMSE

	Linear regression
Original features	2.1608
Normalized features	2.1631

Why might normalizing features in this way not be as helpful as it is for KNN?

Because KNN uses distances, it is sensitive to the scale of features. Linear regression can learn weights to scale features so it is less sensitive. This normalization essentially centers/rescales features.

#### b. Feature selection results

Feature Rank	Feature number	City	Day
1	334	Chicago	-1
2	347	Minneapolis	-1
3	405	Grand Rapids	-1
4	366	Kansas City	-1
5	361	Cleveland	-1
6	307	Omaha	-2
7	367	Indianapolis	-1
8	264	Minneapolis	-2
9	9	Boston	-5
10	236	Springfield	-3

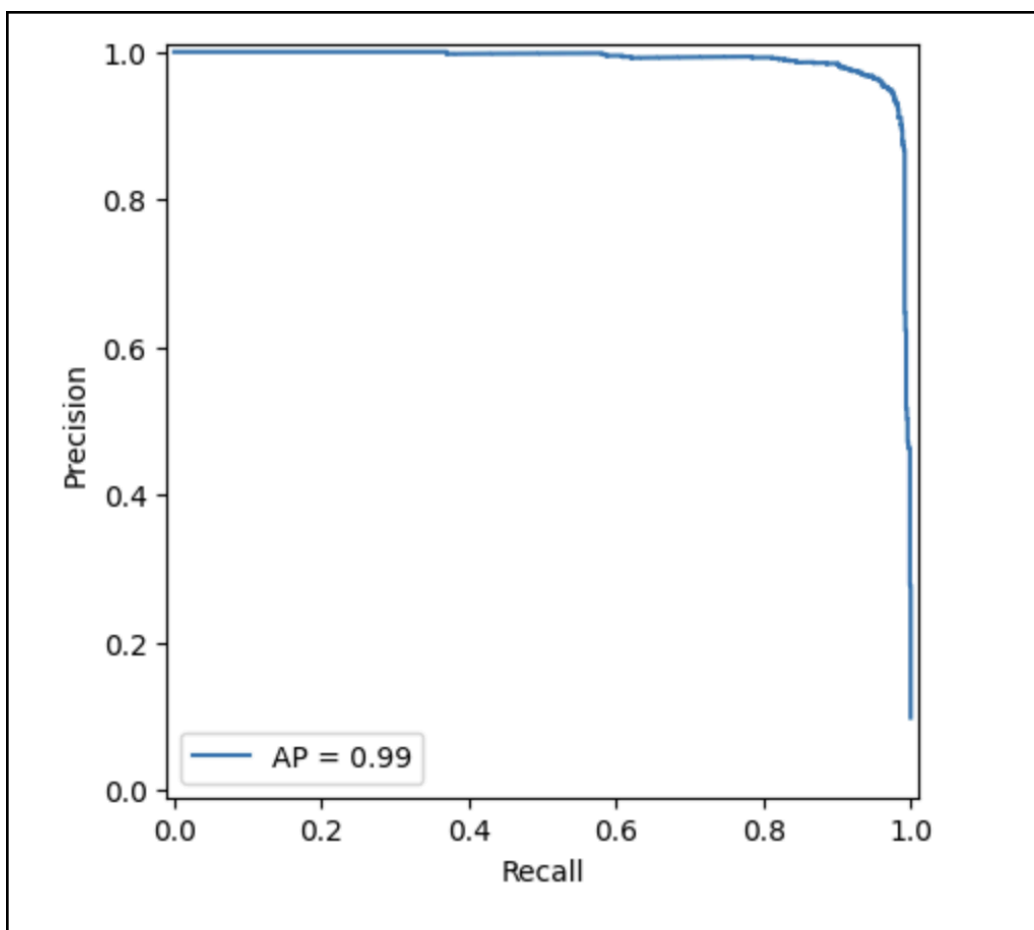
Test error using only the 10 most important features for regression

	Linear Regression
RMS Error	2.0621

#### 4. Stretch Goals

##### a. PR and ROC curves

PR plot

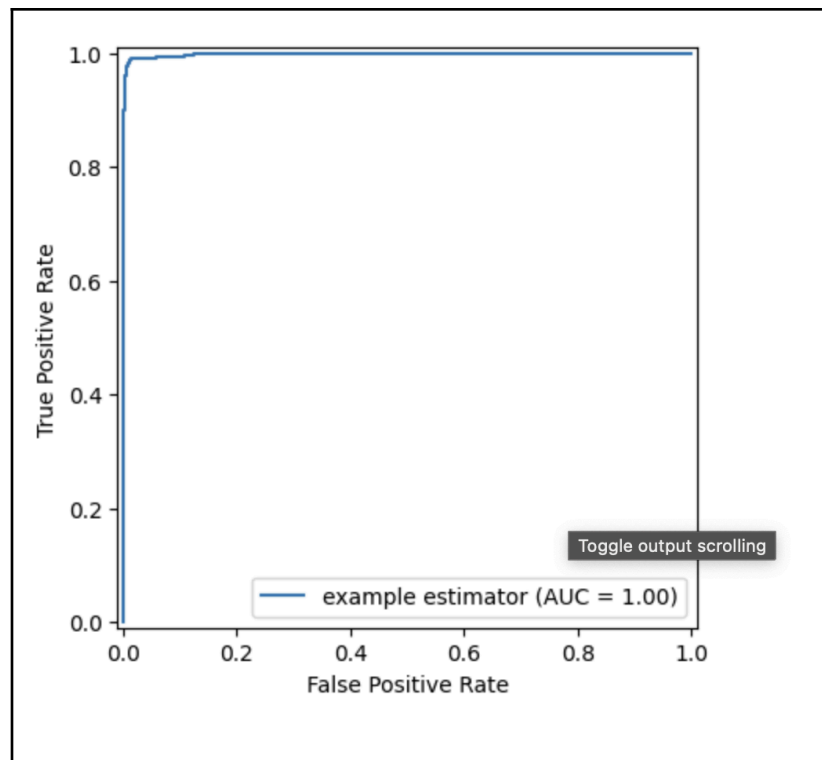


Average Precision

0.98883

ROC plot



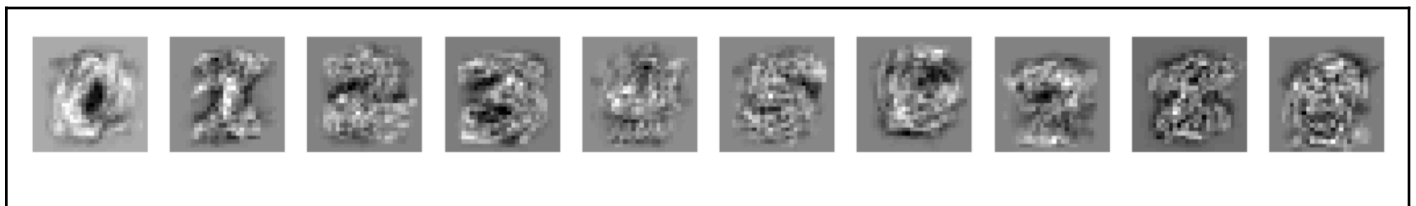


Area under the curve (AUC)


0.9983

**b. Visualize weights**

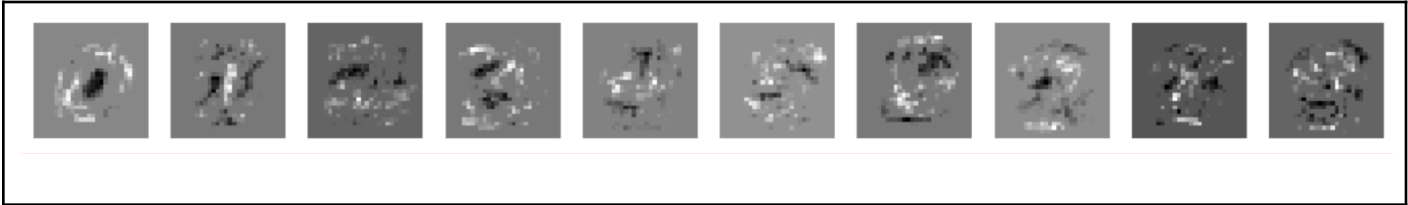
LLR - L2




LLR - L1



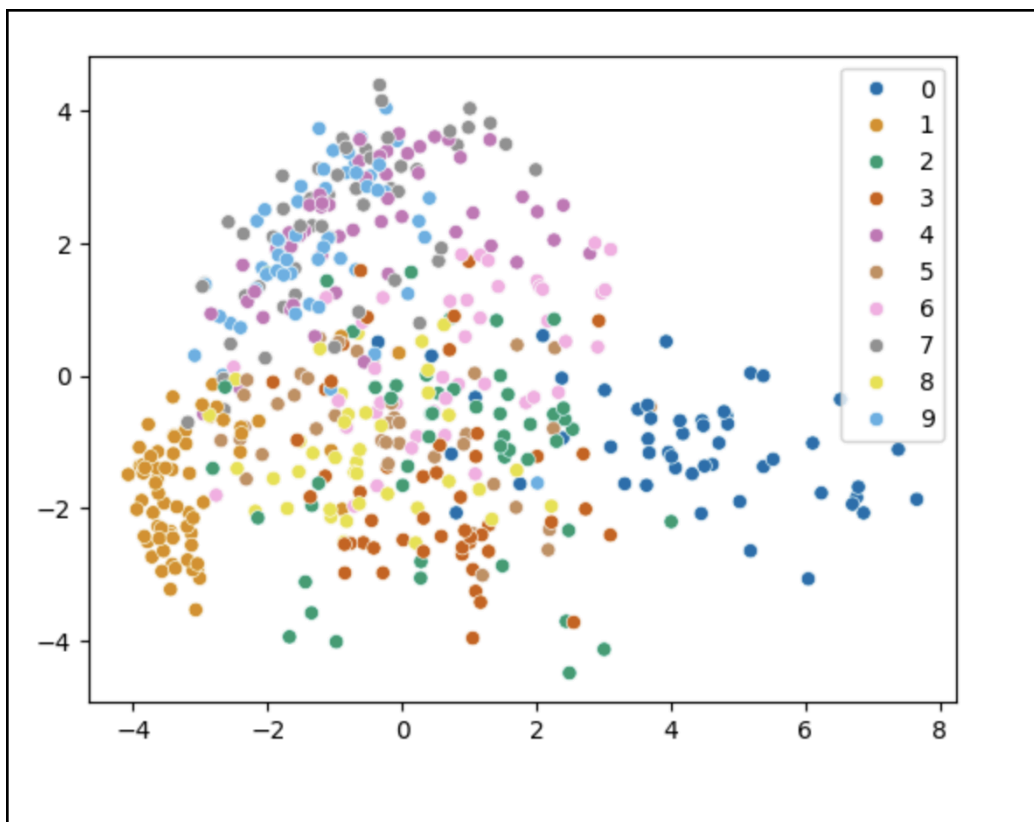
LLR - elastic



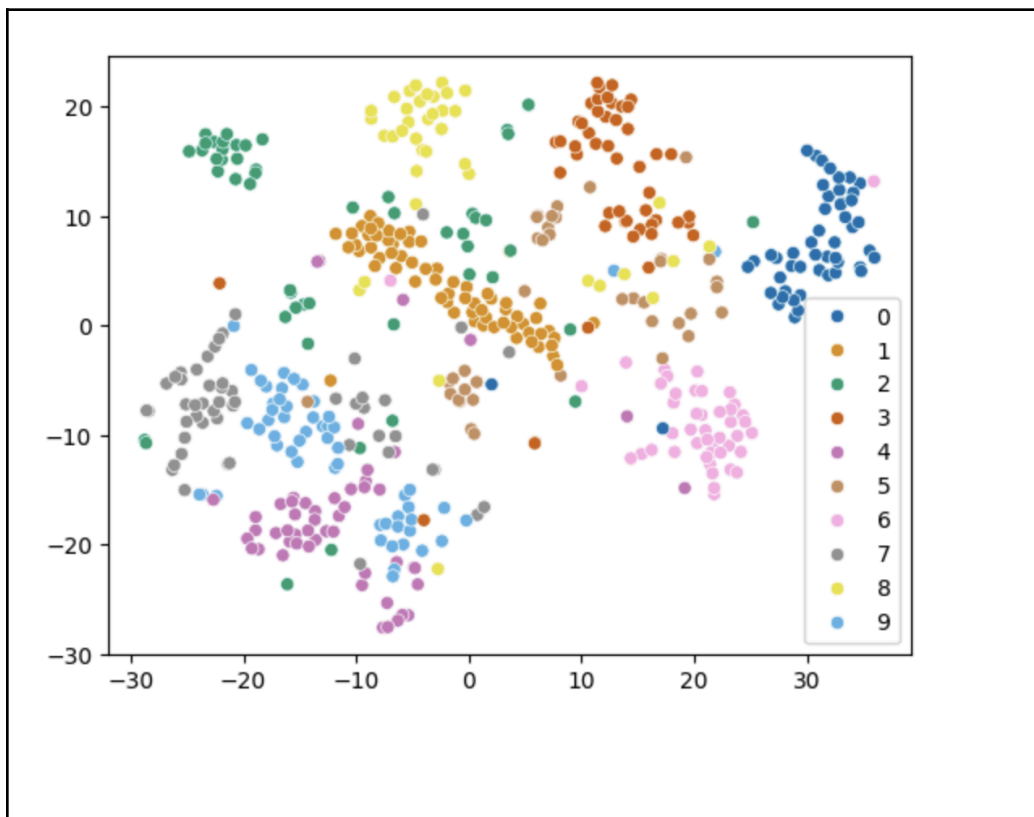
SVM



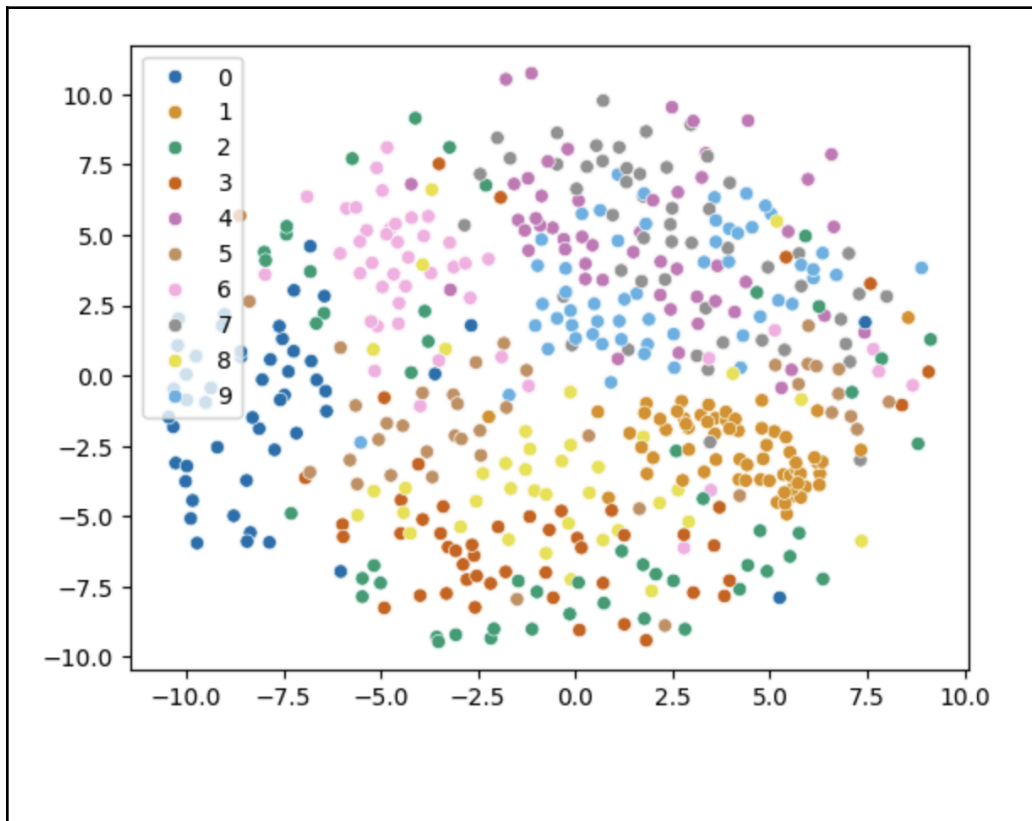
**c. Other embeddings**  
Display 2+ plots for TSNE, MDA, and/or LDA, and copy PCA plot from 1b here.  
PCA



TNSE



MDS



**d. One city is all you need**

City

St. Louis

Test error using features only from that city

3.1263 RMSE

Explain your process (in words):

Iterated through every city, and trained a Ridge model (no feature normalization, default parameters) using only days from that city as input features. Then I took the validation error for each of these models, and took the city with the minimum validation error.

#### e. Compare linear SVM and SVM with RBF kernel

Test accuracy (%)

# training samples	SVM-Linear	SVM-RBF
100	32.4%	34.4%
1,000	16.11 %	9.17%
10,000	11.12 %	4.06%
60,000	8.17 %	2.08%

#### Acknowledgments / Attribution

List any outside sources for code or ideas or "None".  
Stackoverflow + sklearn docs

# CS441\_SP24\_HW2\_Starter

February 20, 2024

## 0.1 CS441: Applied ML - HW 2

### 0.1.1 Parts 1-2: MNIST

Include all the code for generating MNIST results below

```
[253]: # initialization code
import numpy as np
from keras.datasets import mnist
%matplotlib inline
from matplotlib import pyplot as plt
from scipy import stats
from sklearn.linear_model import LogisticRegression

!apt install libomp-dev > /dev/null 2>&1
!pip install faiss-cpu > /dev/null 2>&1
import faiss
import time

def load_mnist():
    """
    Loads, reshapes, and normalizes the data
    """
    (x_train, y_train), (x_test, y_test) = mnist.load_data() # loads MNIST data
    x_train = np.reshape(x_train, (len(x_train), 28*28)) # reformat to 768-d
    ↪vectors
    x_test = np.reshape(x_test, (len(x_test), 28*28))
    maxval = x_train.max()
    x_train = x_train/maxval # normalize values to range from 0 to 1
    x_test = x_test/maxval
    return (x_train, y_train), (x_test, y_test)

def display_mnist(x, subplot_rows=1, subplot_cols=1):
    """
    Displays one or more examples in a row or a grid
    """
    if subplot_rows>1 or subplot_cols>1:
        fig, ax = plt.subplots(subplot_rows, subplot_cols, figsize=(15,15))
        for i in np.arange(len(x)):
```

```

    ax[i].imshow(np.reshape(x[i], (28,28)), cmap='gray')
    ax[i].axis('off')
else:
    plt.imshow(np.reshape(x, (28,28)), cmap='gray')
    plt.axis('off')
plt.show()

```

### 0.1.2 Part 1: PCA and Data Compression

```


[254]: from sklearn.decomposition import PCA
import matplotlib.pyplot as plt

(x_train, y_train), (x_test, y_test) = load_mnist()

# Compute the first 10 principal components using x_train
pca_model = PCA()
pca_model.fit(x_train)
principal_components = pca_model.components_[:10]

# Display First 10 Components
display_mnist(principal_components, 1, 10)

```




```

[57]: # Scatter plot of first two PCA dimensions
import seaborn as sns

# use pca.transform
x = pca_model.transform(x_train[:500])
# TO DO
ind = np.arange(500)
sns.scatterplot(x=x[ind,0],y=x[ind,1], hue=y_train[ind], palette="colorblind")

```

[57]: <Axes: >




```
[58]: # Plot cumulative explained variance ratio
      # cumsum and pca.explained_variance_ratio_ will be useful

      # TO DO
      cum_var = np.cumsum(pca_model.explained_variance_ratio_)
      plt.plot(cum_var)
```

```
[58]: [<matplotlib.lines.Line2D at 0x7f3a416a2190>]
```





```
[59]: # Select number of dimensions that explains 90% of variance, according to your
      ↪ plot above
      m = (cum_var >= 0.9).argmax() + 1
      print(m)

      pca_model = PCA(n_components=m).fit(x_train)

      # Get time and error when using original features with brute force 1-NN
      # TO DO
      start = time.time()
      index = faiss.IndexFlatL2(x_train.shape[1])
      index.add(x_train)
      dist, idx = index.search(x_test, 1)
      pred = y_train[idx.squeeze()]
      err = (pred != y_test).mean()
      print(err)
      print(time.time() - start)

      # Get time and error when using compressed features with brute force 1-NN
      # TO DO
      x_train_comp = pca_model.transform(x_train)
      x_test_comp = pca_model.transform(x_test)
```

```

start = time.time()
index = faiss.IndexFlatL2(x_train_comp.shape[1])
index.add(x_train_comp)
dist, idx = index.search(x_test_comp,1)
pred = y_train[idx.squeeze()]
err = (pred != y_test).mean()
print(err)
print(time.time() - start)

```

```

87
0.0309
36.031790018081665
0.0273
4.545729160308838

```

### 0.1.3 Part 2: MNIST Classification with Linear Models

```

[6]: from sklearn.linear_model import LogisticRegression
     from sklearn.svm import LinearSVC

```

LLR/SVM vs training size

```

[14]: # LLR
      for n in [100, 1000, 10000, 60000]:
          LLR = LogisticRegression(max_iter = 10000).fit(x_train[:n],y_train[:n])
          pred = LLR.predict(x_test)
          err = 100*(pred != y_test).mean()
          print(err)

```

```

32.5
13.639999999999999
9.5
7.3800000000000001

```

```

[7]: # SVM
      for n in [100, 1000, 10000, 60000]:
          SVM = LinearSVC(max_iter = 10000).fit(x_train[:n],y_train[:n])
          pred = SVM.predict(x_test)
          err = 100*(pred != y_test).mean()
          print(err)

```

```

/home/matty/.local/lib/python3.11/site-packages/sklearn/svm/_classes.py:31:
FutureWarning: The default value of `dual` will change from `True` to `auto`
in 1.5. Set the value of `dual` explicitly to suppress the warning.
  warnings.warn(
/home/matty/.local/lib/python3.11/site-packages/sklearn/svm/_classes.py:31:
FutureWarning: The default value of `dual` will change from `True` to `auto`

```

in 1.5. Set the value of `dual` explicitly to suppress the warning.

```
warnings.warn(
```

32.35

16.11

/home/matty/.local/lib/python3.11/site-packages/sklearn/svm/\_classes.py:31:

FutureWarning: The default value of `dual` will change from `True` to `'auto'`

in 1.5. Set the value of `dual` explicitly to suppress the warning.

```
warnings.warn(
```

11.12

/home/matty/.local/lib/python3.11/site-packages/sklearn/svm/\_classes.py:31:

FutureWarning: The default value of `dual` will change from `True` to `'auto'`

in 1.5. Set the value of `dual` explicitly to suppress the warning.

```
warnings.warn(
```

8.17

Error visualization

```
[54]: # to get scores for logistic regression use: scores = model_lr.  
      ↪predict_proba(x_test)  
def makeplots(scores):  
    mostcorrect = np.empty((10, x_test.shape[1]))  
    mostwrong = mostcorrect.copy()  
    for i in range(10):  
        print(f"idx {i}")  
        probs = scores[:,i]  
        islabel = (y_test == i).astype(int)  
  
        maxidx = (probs * islabel).argmax()  
        minidx = ((1-probs) * islabel).argmax()  
        print(probs[maxidx], probs[minidx])  
        # display_mnist(x_test[minidx])  
        # display_mnist(x_test[maxidx])  
        mostcorrect[i] = x_test[maxidx]  
        mostwrong[i] = x_test[minidx]  
  
    display_mnist(mostcorrect, subplot_cols=10)  
    display_mnist(mostwrong, subplot_cols=10)  
  
    makeplots(LLR.predict_proba(x_test))  
    # to get scores for SVM use: scores = model_svm.decision_function(x_test)  
    makeplots(SVM.decision_function(x_test))
```

idx 0

0.9999999963899942 0.00025819211990930365

idx 1

0.9996567977883941 3.0877626333995e-07

idx 2  
 0.9999999744316168 8.143784266293726e-08  
 idx 3  
 0.9999989329684581 2.559199561611065e-05  
 idx 4  
 0.9999997898596669 1.332051810719343e-07  
 idx 5  
 0.9999795469879909 0.00011521805519352722  
 idx 6  
 0.9999992170012151 9.196212267621756e-05  
 idx 7  
 0.9999978520581978 2.8968218320809687e-06  
 idx 8  
 0.9999939555425639 2.7513009347682197e-05  
 idx 9  
 0.9999710313820391 1.0870533123166027e-07



idx 0  
 4.61895263661966 -2.6070372181055292  
 idx 1  
 2.4273329555515355 -5.02282639157594  
 idx 2  
 4.3457220804306536 -4.503737615673133  
 idx 3  
 4.294225965182193 -4.2211613082993695  
 idx 4  
 3.769328884272958 -4.4923061471241255  
 idx 5  
 4.0326785525818565 -4.385248350034097  
 idx 6  
 3.9589625796426087 -3.189884761323814  
 idx 7  
 5.565132589995729 -4.641314083206345  
 idx 8  
 2.6371411849356683 -2.8530038389350647

idx 9

3.0858680559435063 -4.077201020006369



Parameter selection

```
[255]: cs = []  
      errs = []
```

```
[264]: # Try multiple C parameters, select one that minimizes validation error  
      # Often, you need to try a few values and see those results to determine what  
      # other values to try
```

```
def geterr(c):  
    model = LinearSVC(C=c, dual='auto').fit(x_train[:1000], y_train[:1000])  
    pred = model.predict(x_train[50000:])  
    err = (pred != y_train[50000:]).mean()  
    print(c, err)  
    cs.append(c)  
    errs.append(err)  
    return err
```

```
minval = 1000
```

```
minc = -1
```

```
for c in np.linspace(0.0185,0.0199,10):  
    err = geterr(c)  
    if err < minval:  
        minval = err  
        minc = c
```

```
print(minc, minval)
```

0.0185 0.1239

0.018655555555555554 0.1238

0.018811111111111112 0.1237

```

0.018966666666666666 0.1236
0.019122222222222222 0.1237
0.019277777777777778 0.1238
0.019433333333333334 0.1237
0.019588888888888889 0.1235
0.019744444444444446 0.1237
0.0199 0.1238
0.019588888888888889 0.1235

```

```


[265]: # Get test result for selected parameter
plt.semilogx(cs,errs, marker='o', linestyle='')
# TO DO
c = cs[np.argmin(errs)]
print(min(errs))
model = LinearSVC(C=c, dual='auto').fit(x_train[:1000], y_train[:1000])
pred = model.predict(x_test)
err = (pred != y_test).mean()
print(err)

```

```

0.1235
0.1358

```



## 0.2 Part 3: Temperature Regression

```
[125]: import numpy as np
# from google.colab import drive
%matplotlib inline
from matplotlib import pyplot as plt
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso

# load data (modify to match your data directory or comment)
def load_temp_data():
    # drive.mount('/content/drive')
    # datadir = "/content/drive/My Drive/CS441/24SP/hw1/"
    datadir = './'
    T = np.load(datadir + 'temperature_data.npz')
    x_train, y_train, x_val, y_val, x_test, y_test, dates_train, dates_val,
    ↪ dates_test, feature_to_city, feature_to_day = \
    T['x_train'], T['y_train'], T['x_val'], T['y_val'], T['x_test'], T['y_test'],
    ↪ T['dates_train'], T['dates_val'], T['dates_test'], T['feature_to_city'],
    ↪ T['feature_to_day']
    return (x_train, y_train, x_val, y_val, x_test, y_test, dates_train,
    ↪ dates_val, dates_test, feature_to_city, feature_to_day)

# plot one data point for listed cities and target date
def plot_temps(x, y, cities, feature_to_city, feature_to_day, target_date):
    nc = len(cities)
    ndays = 5
    xplot = np.array([-5,-4,-3,-2,-1])
    yplot = np.zeros((nc,ndays))
    for f in np.arange(len(x)):
        for c in np.arange(nc):
            if cities[c]==feature_to_city[f]:
                yplot[feature_to_day[f]+ndays,c] = x[f]
    plt.plot(xplot,yplot)
    plt.legend(cities)
    plt.plot(0, y, 'b*', markersize=10)
    plt.title('Predict Temp for Cleveland on ' + target_date)
    plt.xlabel('Day')
    plt.ylabel('Avg Temp (C)')
    plt.show()
```

```
[126]: # load data
(x_train, y_train, x_val, y_val, x_test, y_test, dates_train, dates_val,
    ↪ dates_test, feature_to_city, feature_to_day) = load_temp_data()
''' Data format:
    x_train, y_train: features and target value for each training sample
    ↪ (used to fit model)
```

```

    x_val, y_val: features and target value for each validation sample (used
    ↪to select hyperparameters, such as regularization and K)
    x_test, y_test: features and target value for each test sample (used to
    ↪evaluate final performance)
    dates_xxx: date of the target value for the corresponding sample
    feature_to_city: maps from a feature number to the city
    feature_to_day: maps from a feature number to a day relative to the
    ↪target value, e.g. -2 means two days before
    Note: 361 is the temperature of Cleveland on the previous day
'''
f = 361
print('Feature {}: city = {}, day= {}'.format(f, feature_to_city[f],
    ↪feature_to_day[f]))
baseline_rmse = np.sqrt(np.mean((y_val[1:]-y_val[:-1])**2)) # root mean squared
    ↪error
print('Baseline - prediction using previous day: RMSE={}'.format(baseline_rmse))


# plot first two x/y for val
plot_temps(x_val[0], y_val[0], ['Cleveland', 'New York', 'Chicago', 'Denver',
    ↪'St. Louis'], feature_to_city, feature_to_day, dates_val[0])
plot_temps(x_val[1], y_val[1], ['Cleveland', 'New York', 'Chicago', 'Denver',
    ↪'St. Louis'], feature_to_city, feature_to_day, dates_val[1])


```

Feature 361: city = Cleveland, day= -1

Baseline - prediction using previous day: RMSE=3.460601246750482







Linear regression test

```
[127]: def normalize_features(x, y, fnum):
    ''' Normalize the features in x and y.
        For each data sample i:
            x2[i] = x[i]-x[i,fnum]
            y2[i] = y[i]-x[i,fnum]
    '''
    x2 = x.copy()
    y2 = y.copy()
    for i in np.arange(len(x)):
        x2[i] = x[i] - x[i, fnum]
        y2[i] = y[i] - x[i,fnum]
    return x2, y2
```

```
[135]: # linear regression (use Ridge)

# original features
model = Ridge().fit(x_train,y_train)
pred = model.predict(x_test)
```

```

print(np.sqrt(np.square(pred-y_test).mean()))

# normalized features
fnum=361
xnorm_train,ynorm_train = normalize_features(x_train,y_train,fnum)
xnorm_test, ynorm_test = normalize_features(x_test,y_test,fnum)

y_pred = Ridge().fit(xnorm_train,ynorm_train).predict(xnorm_test)

print(np.square(y_pred - ynorm_test).mean())**0.5)
# TO DO

```

2.1608605260795755

2.1630698027575606

Feature selection

```

[169]: # feature analysis (select important features using Lasso)
# TO DO
model = Lasso().fit(x_train,y_train)
fidx = (np.abs(model.coef_)>0.001)
fidx = np.arange(x_train.shape[1])[fidx]

import pandas as pd
df = pd.DataFrame({
    'coef': np.abs(model.coef_[fidx]),
    'fnum':fidx,
    'city':feature_to_city[fidx],
    'day':feature_to_day[fidx],
})
df = df.sort_values('coef',ascending=False)
df

```

```

[169]:
   coef  fnum      city  day
6  0.279599  334    Chicago  -1
7  0.225942  347  Minneapolis  -1
12 0.180130  405  Grand Rapids  -1
10 0.162950  366  Kansas City  -1
9  0.138357  361   Cleveland  -1
5  0.111248  307     Omaha  -2
11 0.046103  367  Indianapolis  -1
3  0.043280  264  Minneapolis  -2
0  0.025518    9     Boston  -5
2  0.025198  236  Springfield  -3
4  0.015140  290  Providence  -2
1  0.005399  175     Boston  -3

```

8 0.001785 354 St. Louis -1

```
[267]: # predict using best features
idxs = df['fnum'][:10].values
model = Ridge().fit(x_train[:,idxs],y_train)
pred = model.predict(x_test[:,idxs])
print(np.sqrt(np.square(pred-y_test).mean()))
```

2.7350935573443653

### 0.3 Part 4: Stretch Goals

Include all your code used for any stretch goals in this section. Add headings where appropriate.

#### 1 4a

```
[190]: from sklearn.metrics import precision_recall_curve, roc_curve,
        average_precision_score, auc, roc_auc_score, PrecisionRecallDisplay
import sklearn.metrics as metrics

(x_train, y_train), (x_test, y_test) = load_mnist()
y_train = y_train == 0
y_test = y_test == 0
model = LogisticRegression().fit(x_train,y_train)
```

```
[192]: probs = model.predict_proba(x_test)[:,-1]

average_precision = average_precision_score(y_test, probs)
print('Average Precision (AP):', average_precision)


auroc = roc_auc_score(y_test,probs)
print('Area under the ROC curve (AuROC):', auroc)


precision, recall, _ = precision_recall_curve(y_test, probs)
disp = PrecisionRecallDisplay(precision=precision, recall=recall,
        average_precision=average_precision)
disp.plot()

precision, recall, _ = precision_recall_curve(y_test, probs)
fpr, tpr, thresholds = metrics.roc_curve(y_test, probs)
auroc = roc_auc_score(y_test,probs)
display = metrics.RocCurveDisplay(fpr=fpr, tpr=tpr,
        roc_auc=auroc,estimator_name='example estimator')
display.plot()
fpr, tpr, _ = roc_curve(y_test, probs)
```

Average Precision (AP): 0.9888354276370754

Area under the ROC curve (AuROC): 0.9983122539481425







## 2 4b

```
[212]: (x_train, y_train), (x_test, y_test) = load_mnist()


def go(model):
    model.fit(x_train[:1000], y_train[:1000])
    display_mnist(model.coef_, subplot_cols=10)

go(LogisticRegression(penalty='l1', solver='liblinear'))
go(LogisticRegression(penalty='l2', solver='liblinear'))
go(LogisticRegression(penalty='elasticnet', solver='saga', l1_ratio=0.5))
go(LinearSVC())
```






```
/home/matty/.local/lib/python3.11/site-
packages/sklearn/linear_model/_sag.py:350: ConvergenceWarning: The max_iter was
reached which means the coef_ did not converge
warnings.warn(
```



```
/home/matty/.local/lib/python3.11/site-packages/sklearn/svm/_classes.py:31:
FutureWarning: The default value of `dual` will change from `True` to `auto`
in 1.5. Set the value of `dual` explicitly to suppress the warning.
warnings.warn(
```



### 3 4c

```
[218]: from sklearn.manifold import TSNE
from sklearn.manifold import MDS
(x_train, y_train), (x_test, y_test) = load_mnist()


tsne = TSNE(n_components=2)
X_tsne = tsne.fit_transform(x_train[:500])

# MDS
mds = MDS(n_components=2)
X_mds = mds.fit_transform(x_train[:500])

# Plotting

sns.scatterplot(x=X_tsne[:, 0], y=X_tsne[:, 1], hue=y_train[:500],
↪ palette="colorblind")
```


[218]: <Axes: >



```
[217]: sns.scatterplot(x=X_mds[:, 0], y=X_mds[:, 1], hue=y_train[:500],  
    ↪ palette="colorblind")
```

[217]: <Axes: >





## 4 4D

```
[246]: # load data
(x_train, y_train, x_val, y_val, x_test, y_test, dates_train, dates_val,
  ↪ dates_test, feature_to_city, feature_to_day) = load_temp_data()

def RMSE(a,b):
    return np.sqrt(np.square(a-b).mean())

errs = []
for i in range(83):
    x = x_train[:, i:415:83]
    xv = x_val[:, i:415:83]
    model = Ridge().fit(x, y_train)
    pred = model.predict(xv)
    err = RMSE(pred, y_val)
    errs.append(err)

imin = np.argmin(errs)
print(imin, feature_to_city[imin])
x = x_train[:, imin:415:83]
```

```

model = Ridge().fit(x, y_train)
pred = model.predict(x_test[:,imin:415:83])
err = RMSE(pred, y_test)
print(err)

```

22 St. Louis  
3.1263272970544764

## 5 4e

```

[220]: from sklearn.svm import SVC
# SVM
for n in [100, 1000, 10000, 60000]:
    SVM = SVC(max_iter = 10000).fit(x_train[:n], y_train[:n])
    pred = SVM.predict(x_test)
    err = 100*(pred != y_test).mean()
    print(err)

```

34.410000000000004  
9.17  
4.06  
2.08

```

[ ]: # from https://gist.github.com/jonathanagustin/b67b97ef12c53a8dec27b343dca4abba
# install can take a minute

import os
# @title Convert Notebook to PDF. Save Notebook to given directory
NOTEBOOKS_DIR = "/content/drive/My Drive/CS441/24SP/hw2" # @param {type:
↳ "string"}
NOTEBOOK_NAME = "CS441_SP24_HW2_Solution.ipynb" # @param {type:"string"}
#-----#
from google.colab import drive
drive.mount("/content/drive/", force_remount=True)
NOTEBOOK_PATH = f"{NOTEBOOKS_DIR}/{NOTEBOOK_NAME}"
assert os.path.exists(NOTEBOOK_PATH), f"NOTEBOOK NOT FOUND: {NOTEBOOK_PATH}"
!apt install -y texlive-xetex texlive-fonts-recommended texlive-plain-generic >
↳ /dev/null 2>&1
!jupyter nbconvert "$NOTEBOOK_PATH" --to pdf > /dev/null 2>&1
NOTEBOOK_PDF = NOTEBOOK_PATH.rsplit('.', 1)[0] + '.pdf'
assert os.path.exists(NOTEBOOK_PDF), f"ERROR MAKING PDF: {NOTEBOOK_PDF}"
print(f"PDF CREATED: {NOTEBOOK_PDF}")

```

Mounted at /content/drive/  
PDF CREATED: /content/drive/My Drive/CS441/24SP/hw1/CS441\_SP24\_HW1\_Solution.pdf